MAISON

PRD and Agile Backlog

Your Personal AI Fashion House — From Inspiration to Outfit

> *This PRD builds on the [Product Discovery & Strategy] document. It scopes the MVP defined there — the Wizard-of-Oz-validated "recreate this look within my budget" flow — into requirements ready for design and engineering.*

---

1. Product Objective

The primary objective of MAISON is to reduce the time and effort between a user discovering an outfit they love and being able to purchase a complete, budget-appropriate version of it — regardless of which retailers the individual pieces come from. By replacing manual, cross-platform searching with curated, context-aware outfit recommendations, MAISON helps users move from inspiration to purchase decision with confidence instead of decision fatigue.

---

2. Problem and Goal

2.1 Problem Being Solved

E-commerce platforms are optimized for individual product discovery, while users think in complete outfits, occasions, and personal style. This forces users to manually discover, compare, style, and purchase complementary items across five or more platforms (Pinterest, Google Lens, Myntra, Amazon, Ajio, H&M) to assemble a single outfit — a process that produces decision fatigue rather than product scarcity. Existing retailers (Myntra, Savana, Newme) cannot solve this because their business incentive is to sell their own catalog, not to find the objectively best outfit across the market.

2.2 User Goal

To turn a piece of fashion inspiration — a screenshot, a Pinterest save, or a plain-language description — into 2–3 complete, trustworthy outfit options that fit a stated budget and occasion, without opening five different shopping apps.

2.3 Business Goal

Validate strong desirability and activation within the MVP audience (urban women, 20–30, Pinterest/Instagram-driven shoppers) by achieving a high Outfit Completion Rate and Decision Confidence Score, establishing MAISON as a trusted, retailer-agnostic styling layer that can later scale into affiliate commerce, wardrobe intelligence, and B2B retail partnerships.

2.4 Product Goal

Build a zero-setup, conversational outfit-generation workflow that consumes an inspiration image or a plain-language prompt (occasion + budget) and returns 2–3 complete, purchasable outfit options — sourced across multiple retailers — within 15 seconds, with an interim acknowledgment shown to the user within 2 seconds to preserve the feeling of a responsive, conversational interaction.

---

3. User Stories

User Story 1

Inspiration-to-Outfit Recreation

As an inspiration-driven shopper, I want to upload a screenshot of an outfit I found on Pinterest or Instagram, so that I can receive similar, purchasable pieces without manually searching for every item myself.

User Story 2

Text-to-Outfit Prompt

As a busy professional without a reference image, I want to describe my occasion and budget in plain language (e.g., “business casual outfit under ₹5,000 for a client meeting”), so that I can get a complete outfit without needing to find or upload inspiration first.

User Story 3

Budget-Constrained Outfit Generation

As a value-conscious shopper, I want every recommended outfit to respect a budget I set, so that I don’t have to manually check and add up prices across multiple product pages.

User Story 4

Multi-Retailer Comparison & Buy Links

As a shopper who doesn’t want to be limited to one retailer’s catalog, I want each recommended outfit to show which retailer each piece comes from and a direct buy link, so that I can act on a recommendation immediately without further searching.

User Story 5

Recommendation Confidence Feedback

As a user reviewing curated outfits, I want a fast way to rate how confident I feel about a recommendation, so that MAISON can learn what “a good outfit” means for me and improve future suggestions.

User Story 6

Alternative Suggestion on Size/Stock Unavailability

As a shopper who has found the right outfit, I want MAISON to immediately suggest a visually similar, in-stock alternative if an item is unavailable in my size, so that I don’t have to restart my search from scratch.

---

4. Acceptance Criteria

User Story 1: Inspiration-to-Outfit Recreation

Scenario: Successful processing of an uploaded inspiration screenshot.

Given the user is on the outfit capture screen **When** the user uploads a screenshot of an outfit and sets a budget of ₹5,000 **Then** the system must isolate the distinct garments and accessories in the image, source visually similar in-stock alternatives across at least two retailers, and return 2–3 complete outfit options that collectively stay within the stated budget, within 15 seconds.

User Story 2: Text-to-Outfit Prompt

Scenario: Successful outfit generation from a plain-language prompt with no image.

Given the user is on the outfit capture screen and has not uploaded an image **When** the user types “business casual outfit under ₹5,000 for a client meeting” and submits **Then** the system must extract the occasion (“client meeting”), the style register (“business casual”), and the budget (“₹5,000”), and return 2–3 complete outfit options appropriate to all three constraints within 15 seconds.

User Story 3: Budget-Constrained Outfit Generation

Scenario: Outfit generation correctly respects a hard budget ceiling.

Given the user has set a budget of ₹4,000 **When** the system generates outfit recommendations **Then** the total price of every recommended outfit (sum of all included items) must be less than or equal to ₹4,000, and each outfit must display its total price prominently alongside the individual item prices.

User Story 4: Multi-Retailer Comparison & Buy Links

Scenario: A recommended outfit includes items from more than one retailer.

Given the system has generated an outfit consisting of a top from Retailer A and shoes from Retailer B **When** the outfit is displayed to the user **Then** each item must show its source retailer name and a working buy link that opens the correct product page, and the outfit view must make clear that items come from different retailers rather than presenting them as a single-retailer bundle.

User Story 5: Recommendation Confidence Feedback

Scenario: User rates a recommended outfit after viewing it.

Given the user has viewed a complete outfit recommendation **When** the user is prompted with "How confident do you feel about this outfit?" and selects a rating from 1–5 **Then** the system must record the rating against that specific outfit and user session within 1 second, without requiring the user to leave the current screen or complete any additional steps.

User Story 6: Alternative Suggestion on Size/Stock Unavailability

Scenario: A recommended item is out of stock in the user's selected size.

Given the user has selected their size and one item in a recommended outfit is unavailable in that size **When** the outfit is displayed **Then** the system must automatically replace that item with a visually similar, same-price-tier alternative that is available in the user's size, and must not display out-of-stock items as if they were purchasable.

---

5. Functional Requirements

1. Input & Context Capture

- ✓ The system must accept image uploads (screenshot or photo) in standard formats (JPEG, PNG, WebP) up to 10MB.
- ✓ The system must accept unstructured text prompts up to 500 characters describing occasion, style, and/or budget.
- ✓ The system must provide explicit input fields for occasion and budget that can be set independently of, or in combination with, an image or text prompt.
- ✓ The system must support an optional user-selected size profile to filter recommendations for availability.

2. AI Outfit Generation Engine

- ✓ The system must parse image and/or text input to identify garment categories, style attributes, occasion, and budget constraints.
- ✓ The system must return a structured output of 2–3 complete outfits, each composed of complementary, coordinated items (not isolated single products).
- ✓ The system must ensure every returned outfit's total cost is at or below the user's stated budget.
- ✓ The system must source candidate items from a product catalog spanning multiple retailer partners, not a single retailer's inventory.
- ✓ The AI engine must apply a stylist-like persona: presenting curated options and explaining *why* an outfit fits the brief, without being prescriptive about the user's final choice.

3. Multi-Retailer Product Matching

- ✓ The system must maintain a synchronized product feed (via affiliate APIs or partnered retailer integrations) including price, size availability, and stock status.
- ✓ The system must detect out-of-stock or size-unavailable items at recommendation time and substitute a visually similar in-stock

alternative automatically.

- ✓ The system must attribute every item in an outfit to its source retailer and generate a working, trackable buy link.

4. Outfit Presentation & Execution UI

- ✓ The interface must display 2–3 complete outfits side by side (or in a swipeable sequence on mobile), each showing all included items, total price, and source retailers.
- ✓ The UI must provide distinct, immediate action elements per outfit: **[View Buy Links]**, **[Rate This Outfit]**, and **[Show Alternatives]**.
- ✓ The UI must clearly separate "why this outfit" context (occasion, style match) from the transactional details (price, retailer, stock).

5. Feedback & Rating

- ✓ The system must capture a 1–5 confidence rating per outfit recommendation, tied to the specific outfit and session.
- ✓ The system must log which outfit (if any) a user clicked a buy link for, to measure recommendation acceptance separately from stated confidence.

---

6. Non-Functional Requirements

1. **Response Time (Performance):** The system must display an interim acknowledgment (e.g., "Maison is curating your look…") within 2 seconds of submission, and return the complete set of 2–3 outfit recommendations within 15 seconds, to preserve a conversational, non-frustrating experience.
2. **Accessibility (UI/UX):** The app interface must conform to WCAG 2.1 AA guidelines, with high-contrast styling, clear visual hierarchy between outfits, and screen-reader-compatible image descriptions for uploaded and recommended items.
3. **Privacy & Data Protection:** Uploaded images, style preferences, and budget data must be encrypted in transit and at rest, and excluded from upstream LLM training data loops via enterprise API agreements. Users must be able to delete uploaded images and history on request.
4. **Reliability & Fallbacks:** The core recommendation flow must maintain a 99.5% uptime standard. If the AI/product-matching engine experiences an outage or severe latency spike, the system must degrade gracefully to a small set of pre-curated, occasion-based outfit sets rather than failing outright — consistent with the Wizard-of-Oz validation approach used pre-launch.
5. **Catalog Freshness:** Product price, size availability, and stock data pulled from retailer feeds must not be stale by more than 24 hours, to avoid recommending unavailable or mispriced items.
6. **Scalability:** The recommendation engine and product-matching layer must be architected to add new retailer feeds without requiring changes to the core outfit-generation logic.

---

7. Assumptions

1. Users are willing to upload personal screenshots or photos of outfits they did not create themselves (sourced from social media, influencers, or celebrities) for the purpose of recreation.
2. Users trust AI-generated styling recommendations enough to consider clicking through to purchase, rather than treating them purely as inspiration.
3. Users are willing to have their outfit sourced across multiple retailers rather than requiring a single-brand bundle.
4. Budget is a meaningful constraint for users, but not the sole driver of outfit choice — users will accept a well-matched outfit that costs somewhat more over a cheaper, weaker match.
5. A sufficiently broad multi-retailer product feed (via affiliate networks or direct partnerships) is obtainable at MVP stage without requiring exclusive data-sharing agreements from competitor-adjacent retailers.
6. Users shopping via MAISON are open to discovering and buying from retailers outside their usual shopping habits, provided the recommendation is trustworthy.

---

8. Dependencies

1. AI Services

- Availability of image-understanding and language-model providers (e.g., OpenAI, Anthropic) for outfit parsing and generation.
- Consistent response quality and latency from AI APIs to meet the 15-second recommendation target.

2. Retail Data & Commerce

- Access to multi-retailer product feeds with price, size, and stock data — via affiliate networks (e.g., existing fashion affiliate programs) or direct retailer partnerships.
- Trackable, working buy-link generation per retailer for attribution and monetization.

3. Design Resources

- UX design support for the image-upload and text-prompt capture flows.
- Accessibility review for outfit-comparison and recommendation UI.

4. Engineering Resources

- Frontend development for the outfit capture, comparison, and rating interfaces.
- Backend infrastructure for AI orchestration, product-matching, and catalog synchronization.

5. Authentication & User Management

- Support for email-based login, magic links, or social sign-in.
- Secure session management to persist budget, size, and style preferences across visits.

6. Legal & IP Review

- Review of image-upload and outfit-recreation flow for copyright and IP exposure (recreating outfits sourced from third-party or celebrity images).
- Review of affiliate/retailer partnership terms for compliance with each retailer's API usage policies.

7. Analytics & Monitoring

- Product analytics for activation, Outfit Completion Rate, and Decision Confidence Score tracking.
- Monitoring tools for uptime, latency, and product-feed staleness.

---

9. Risks & Mitigations

Risk 1: Monetization Model Conflicts With Core Positioning Category: Business Risk Description: MAISON's value proposition is "the best outfit regardless of retailer," but affiliate-commission monetization creates an incentive to favor higher-commission retailers over the objectively best match — undermining user trust if discovered. Mitigation: Define and disclose a transparent ranking policy; consider a subscription or flat-fee model as an alternative or complement to commission-based revenue; monitor recommendation patterns for commission-driven bias.

Risk 2: Multi-Retailer Catalog Access Is Difficult to Obtain Category: Technical / Business Risk Description: Retailers who view MAISON as competitive (per the Myntra incentive-conflict analysis) have limited motivation to provide clean API access, risking thin or unstable catalog coverage. Mitigation: Start with a small set of partnered retailers and affiliate-network feeds rather than the full open market; expand retailer coverage incrementally as trust and traffic value are demonstrated.

Risk 3: AI Generates Ineffective or Mismatched Outfits Category: Product Risk Description: The AI may generate outfits that are stylistically incoherent, poorly fitted to the stated occasion, or exceed budget due to parsing errors, reducing user trust and recommendation acceptance. Mitigation: Implement prompt engineering guidelines and structured output validation (including a hard budget-ceiling check before display); use the "Rate This Outfit" and "Show Alternatives" feedback loops to continuously evaluate output quality.

Risk 4: Low User Retention Beyond Novelty Category: User Risk Description: Users may find the first few recommendations delightful but stop returning if outfit quality doesn't improve or feel personalized over time. Mitigation: Track Outfit Completion Rate and Decision Confidence Score as leading indicators; introduce lightweight personalization from rating history; expand into occasion-based planning and wardrobe completion per the roadmap once core recreation flow is validated.

Risk 5: AI/Catalog-Matching Latency Category: Technical Risk Description: Multi-retailer product matching adds latency beyond a typical single-model AI call, risking a response time that breaks the conversational feel of the experience. Mitigation: Use an interim "curating your look" acknowledgment within 2 seconds; optimize catalog lookup with pre-indexed embeddings; fall back to pre-curated occasion-based sets if live matching exceeds the latency threshold.

Risk 6: Privacy and Trust Concerns Around Image Uploads Category: User / Business Risk Description: Users may hesitate to upload personal screenshots or describe budget constraints to an AI-powered system, particularly given the sensitivity of financial and body-related preferences. Mitigation: Clearly communicate privacy policy and data retention; encrypt image and preference data; exclude user data from upstream LLM training loops; provide an easy image/history deletion option.

Risk 7: IP/Copyright Exposure From Outfit Recreation Category: Legal Risk Description: Ingesting screenshots of outfits worn by influencers or celebrities to source purchasable alternatives sits in a legally ambiguous space at scale. Mitigation: Legal review of the recreation flow prior to scaling; ensure MAISON sources and recommends *similar* items rather than reproducing or redistributing the original image/content itself.

Risk 8: Weak Problem-Solution Fit at Scale Category: Product Risk Description: The Wizard-of-Oz validation may show demand for human-curated recommendations without proving that an AI + live multi-retailer catalog can replicate that quality automatically. Mitigation: Treat desirability validation and technical feasibility as separate, sequential gates; do not scale retailer integrations until the Wizard-of-Oz signal clears the go/no-go threshold defined in the discovery phase.

---

10. Agile Backlog

Epic 1: Inspiration-to-Outfit Generation

Enable users to turn an inspiration image or plain-language prompt into complete, budget-appropriate outfit recommendations.

Epic	User Story	Priority	Story Points	Notes
Inspiration-to-Outfit Generation	As a user, I want to upload a screenshot of an outfit so that I can receive similar, purchasable items without manual searching.	High	8	Requires image-parsing/garment-detection capability
Inspiration-to-Outfit Generation	As a user, I want to describe my occasion and budget in plain text so that I can get an outfit without needing a reference image.	High	5	Text-prompt parsing (occasion, style, budget extraction)
Inspiration-to-Outfit Generation	As a user, I want to set a budget that every recommended outfit respects so that I don't have to manually total prices.	High	5	Hard budget-ceiling validation before display
Inspiration-to-Outfit Generation	As a user, I want to select an occasion so that recommendations match the context I'm dressing for.	High	3	Occasion taxonomy (interview, wedding, date, office, etc.)
Inspiration-to-Outfit	As a user, I want to	High	8	Core value

Generation	see 2–3 complete outfit options instead of hundreds of individual products so that I can decide quickly.			proposition; depends on AI generation + catalog matching

Epic 2: Multi-Retailer Shopping & Decision Support

Enable users to act on outfit recommendations with confidence, across retailers, without losing progress to unavailable items.

Epic	User Story	Priority	Story Points	Notes
Multi-Retailer Shopping & Decision Support	As a user, I want to see which retailer each item comes from and a direct buy link so that I can purchase immediately.	High	5	Requires affiliate/retailer feed integration
Multi-Retailer Shopping & Decision Support	As a user, I want unavailable items automatically replaced with in-stock alternatives so that I don't have to restart my search.	Medium	5	Depends on live stock/size data freshness
Multi-Retailer Shopping & Decision Support	As a user, I want to rate how confident I feel about a recommended outfit so that future suggestions improve.	Medium	3	Feeds the Decision Confidence Score north star metric
Multi-Retailer Shopping & Decision Support	As a user, I want to request alternatives to a recommended outfit so that I have options if none of the first set feels right.	Medium	5	"Show Alternatives" UI action
Multi-Retailer Shopping & Decision Support	As a user, I want my size and budget preferences to persist across visits so that I don't have to re-enter them each time.	Low	3	Requires authentication/session management

---

This PRD reflects the MVP scope validated through the Wizard-of-Oz desirability experiment described in the Product Discovery & Strategy document. Feasibility work (catalog sourcing, AI matching accuracy) and viability work (monetization model resolution) should progress against the risks listed in Section 9 before full-scale retailer integration begins.