

MAISON — AI Feature and Responsible AI

Your Personal AI Fashion House, From Inspiration to Outfit

July 2026

1. AI Feature Overview

What the AI Feature Does

MAISON’s core AI feature converts an unstructured inspiration image or a plain-language brief into 2–3 complete, coordinated, budget-appropriate outfits sourced across multiple retailers. When a user uploads a screenshot of an outfit they love, or types something like “business casual outfit under ₹5,000 for a client meeting,” the AI isolates the individual garments or extracts the occasion, style register, and budget, then assembles complementary items into full outfits — not a list of individually matched products — and explains, in plain language, why each outfit fits the brief.

Where It Appears in the User Journey

The AI activates on the outfit capture screen, the moment between a user’s raw input (an uploaded image or a typed brief) and the outfit comparison view where 2–3 curated options are displayed side by side. It also re-engages whenever a user taps “Show Alternatives,” rates an outfit, or when an item becomes unavailable in their size — recalibrating the recommendation based on that signal without requiring the user to restart their search.

Why AI Is Essential for This Problem

Rule-based, single-retailer search requires structured input: a category, a size, a price filter. A user who has only a vague visual impression of a look — “something like this, but I don’t know what any of it’s called” — cannot produce that structured input, and that gap is precisely the problem MAISON exists to close. Large models are uniquely suited here because they can interpret an image or an ungrammatical, emotionally-toned brief and extract garment identity, style compatibility, and budget constraints without requiring the user to do the translation work themselves — and because assembling *complementary* items across catalogs that were never designed to talk to each other is a judgment task, not a lookup task.

2. Type of AI Capability Used

MAISON combines several AI capabilities:


- 1. **Computer Vision / Image Understanding** to isolate individual garments and accessories from an uploaded screenshot or photo.
- 2. **Natural Language Understanding** to extract occasion, style register, and budget from unstructured text prompts.
- 3. **Cross-Catalog Matching / Retrieval** using visual and style embeddings to find visually similar, in-stock alternatives across partnered retailer catalogs.
- 4. **Personalization Engine** to weight future recommendations using a user’s rating history and accepted/rejected outfits.
- 5. **Generative Stylist Persona** to explain, in natural language, why a given outfit fits the occasion and budget — without being prescriptive about the user’s final choice.

3. Input and Output Matrix

Input Given to AI	AI Processing / Logic	Output Shown to User
Inspiration screenshot + budget (₹5,000)	Isolates garments via image segmentation; extracts silhouette, color, and style attributes; queries multi-retailer catalog embeddings for visually similar in-stock items; assembles combinations that collectively stay within budget.	“3 outfits inspired by your look — all under ₹5,000, sourced from Myntra, Ajio & H&M.”
Text prompt: “business casual outfit under ₹5,000 for a client meeting”	Extracts occasion, style register, and budget; maps the brief to a style-embedding space; queries the catalog for matching, complementary items.	“Outfit A: Tailored blazer + trousers — ₹4,650 total, from 2 retailers.”
“Rate This Outfit” tap (1-5)	Logs the confidence score against that specific outfit and session; feeds the personalization engine to adjust future style and price-tier weighting.	“Got it — we’ll fine-tune future picks toward what feels right for you.”
Item unavailable in selected size	Detects unavailability at recommendation time; retrieves a visually similar, same-price-tier alternative that is in stock in the user’s size.	“Swapped: this top wasn’t in your size — here’s a similar one from Ajio, same price range.”

4. AI Workflow Step-by-Step

- Ingestion:** The user submits an inspiration image, a text brief, or both, along with an optional size profile.
- Context Assembly:** The system pairs the raw input with active state — occasion, budget, size, and any style-history signal from prior sessions.
- AI Matching Engine (Vision + LLM):** The grouped payload enters the model pipeline, where a vision model isolates garments (if an image was supplied) and an LLM reasons over few-shot examples to assemble complementary items into complete outfit candidates.
- Budget & Safety Filter:** The candidate outfits pass through automated validation — a hard budget-ceiling check, a catalog-freshness check, and a content/IP safety pass — before anything is eligible to render.
- Structure & Retailer Guard:** The validated output is checked for JSON structure, working buy links, and live stock/size status per item, ensuring nothing unpurchasable is ever displayed as available.
- Outfit Card Display:** The interface reveals 2–3 complete outfit cards with total price, source retailers, and a short “why this outfit” explanation.



MAISON AI Workflow

The feedback loop closes back into Context Assembly: a rating, a “Show Alternatives” request, or an automatic out-of-stock swap all update the state payload used for the next interaction, rather than starting the user over from scratch.

5. Data Architecture & Sensitive Information Boundaries

Data Utilized by the System

- User-Provided Operational Data:** Uploaded inspiration images and free-text occasion/budget briefs submitted during an outfit-recreation session.
- Volitional Context Metadata:** Selected occasion, budget, size profile, and style preferences.
- Session Progress Metrics:** Confidence ratings, buy-link clicks, and “Show Alternatives” requests.

Sensitive Data Policy & Compliance

Because uploaded images may depict real people — the user themselves, or third parties such as celebrities and influencers whose looks inspired the request — and because budget and body-adjacent size data are personally sensitive, user inputs are handled under a defined protocol:

- No Biometric or Body-Shape Inference:** The system analyzes images strictly for garment identity, color, and silhouette. It does not perform facial recognition, does not infer body shape or physical attributes, and relies only on the size the user explicitly selects — never an inferred one.
- Feature-Extraction-Only Image Handling:** An uploaded inspiration image is used to extract garment and style embeddings for the duration of the session. It is not redistributed, is not shared with retailer partners, and is not retained beyond what is required to

serve and, if requested, re-serve that session — directly supporting the IP/copyright posture in Section 7.

3. **Privacy-First Exclusions:** All connections to upstream LLM and vision-model API endpoints use enterprise data-privacy agreements, ensuring uploaded images and text are excluded from provider model-training loops. Images and history are deletable by the user on request, with encryption in transit and at rest.

6. Possible AI Failure Cases

1. **Garment Misidentification (Failure to Decompose the Look)** The vision model misreads a garment — e.g., classifying a blazer as a trench coat — producing an outfit that is stylistically incoherent with what the user actually uploaded.
2. **Contextual Misunderstanding of Occasion or Style Language** A stylistic or cultural term is taken too literally or too narrowly — e.g., “business casual” resolved to a Western corporate register when the user’s context implies a different local formalwear norm — producing outfits that technically match the words but not the intent.
3. **Over-Literal Budget Optimization** The system prioritizes staying under budget over style coherence, returning outfits that are technically compliant but visually mismatched — cheapest-fit rather than best-fit.
4. **Catalog/API Latency** Live multi-retailer stock and price lookups push response time beyond the 15-second target, risking abandonment before any outfit is shown.
5. **Hallucinated Availability or Pricing** The AI presents a buy link or price for an item that is actually out of stock or mispriced due to stale catalog data, leading a user to a dead end at the point of purchase.
6. **Cultural or Body-Type Bias in Styling** A model trained predominantly on Western fashion imagery may under-serve Indian body types, skin tones, or occasion norms — e.g., under-recommending ethnic wear for festival occasions where it would be the natural first choice.

7. Responsible AI Risks

1. **Privacy Risk** Users may upload personal photos (not just inspiration screenshots) that incidentally expose their face, location, or other identifiable people in the frame.
2. **IP/Copyright Risk** Ingesting screenshots of outfits worn by celebrities or influencers to source purchasable near-matches sits in a legally ambiguous space, particularly once the flow moves beyond a human-curated validation stage to AI-touching automation.
3. **Bias & Fairness Risk** Style, body-type, and cultural bias in recommendations can alienate users whose context doesn’t match the model’s training distribution.
4. **Hallucination Risk** Fabricated pricing, stock status, or retailer attribution can lead a user to click through to something that isn’t actually purchasable as shown.
5. **Monetization Trust Risk** If retailer commission rates are allowed to vary and influence ranking, the AI’s recommendations could quietly favor higher-commission partners over the objectively best match — directly undermining the retailer-agnostic promise the product is positioned on.
6. **Appearance-Anxiety / Overdependence Risk** Because fashion recommendations are inherently tied to personal appearance and self-image, users could become reliant on AI validation for style confidence, or experience the “Not quite my style” feedback loop as a judgment on their taste rather than a product signal.

8. Guardrails and Mitigations Matrix

Risk	Possible Impact	Mitigation / Guardrail
IP/Copyright exposure	Legal exposure from ingesting and processing third-party or celebrity images at scale.	Feature-extraction-only image handling (no redistribution or long-term retention); legal review of the recreation flow completed and signed off before the first AI-touching (non-Wizard-of-Oz) user, not deferred to a later scaling milestone.
Privacy — identifiable people in images	Accidental exposure of a user’s face or a third party’s likeness via an uploaded photo.	Automatic face-region exclusion before any image reaches garment-detection processing; on-screen guidance nudging users toward inspiration screenshots rather than personal photos.
Hallucinated pricing / availability	User clicks through to an item that is not actually available as shown, damaging trust at the point of purchase.	Hard budget-ceiling and live-stock validation cross-checked against the catalog feed before display; any item failing the 24-hour freshness check is automatically excluded rather than shown stale.
Bias & fairness	Culturally misaligned or body-type-insensitive recommendations alienate Indian users.	Culturally diverse few-shot prompt templates calibrated for Indian occasions and body types; regional style review in prompt design; ongoing bias audits against confidence-rating and “Show Alternatives” signals.
Monetization trust risk	Commission-driven ranking bias undermines the retailer-agnostic promise.	Commission rates normalized across partnered retailers as a condition of partnership; ranking policy published and periodically audited — the direct resolution to the monetization-vs-positioning conflict identified during discovery review.

Appearance-anxiety / overdependence	Users treat AI style judgments as authoritative rather than as suggestions.	All outfit cards framed as suggestions, never instructions (“This could work for your brief” rather than “You should wear this”); binary, low-stakes feedback (“Not quite my style,” not “wrong”); no streak or shame mechanics tied to rejected outfits.
-------------------------------------	---	---

9. AI Transparency & Communication Strategy

- 1. Clear AI Labeling** The outfit comparison view carries a persistent label: “Curated by MAISON’s AI stylist — not a personal stylist or fashion authority.”
- 2. Explicit Onboarding Disclaimer** Before a user’s first session, a mandatory modal discloses that uploaded images are used only to extract style features for that session, are not stored long-term or shared with retailers, and that recommendations are AI-generated suggestions, not professional styling advice.
- 3. Active Feedback Signals** Every outfit card carries “Rate This Outfit” and “Show Alternatives” controls, giving the user direct, low-friction oversight of whether a recommendation landed.
- 4. Confidence Indicators** Where the AI has low contextual confidence — for example, a low-resolution image or a very short text brief — the outfit card surfaces a soft indicator: “We had limited detail to work with here — does this feel right?”
- 5. Ranking Transparency** A visible note on the comparison view discloses the ranking basis in plain language: “Outfits are ranked by fit to your brief, style match, and price — not by which retailer pays us.” This is the user-facing counterpart to the commission-normalization guardrail in Section 8.