

MAISON — Metrics, Experimentation and GTM

Your Personal AI Fashion House, From Inspiration to Outfit

July 2026

1. The North Star Metric

Decision Confidence Score (DCS)

- Definition:** The average self-reported rating (1–5) a user gives immediately after viewing or selecting a recommended outfit, in response to the prompt “How confident do you feel about this outfit?”
- Why It Matters:** Traditional shopping apps optimize for “time spent browsing” or “products viewed.” For a decision-fatigue tool, success is the opposite. We want to minimize the time and doubt a user spends comparing tabs, and maximize how quickly they feel sure of a choice.
- Connection to User Value:** A rising Decision Confidence Score directly correlates with closing the gap between inspiration and purchase. It proves the system replaced manual cross-platform comparison with a recommendation the user actually trusts.
- Connection to Business Value:** High-confidence outfit sessions build the psychological loop of “someone already did the hard part for me.” That loop is what turns a one-off styling favor into a habitual first stop before any purchase — driving retention and the screenshot-sharing behavior already observed in user research (Insight #8: users seek external validation before buying).

Pairing note: Decision Confidence Score is a self-reported leading indicator and is noisy on its own. It is always reported alongside Recommendation Acceptance Rate (buy-link click-through — see Supporting KPIs below) so that a rise in stated confidence unaccompanied by a rise in actual click-through is visible early, not discovered only at a later retention review.

2. Supporting KPIs

Metric	What It Measures	Why It Matters
Outfit Completion Rate	The percentage of sessions where a user views the full set of 2–3 generated outfits rather than abandoning mid-generation.	Measures whether the curation felt trustworthy enough for the user to actually reach a complete look, not just a partial result.
Recommendation Acceptance Rate	The % of recommended outfits where the user clicks at least one buy link.	Tracks real purchase intent as a behavioral counterpart to self-reported confidence — the primary lagging pair to the North Star.
Budget Adherence Accuracy	The % of generated outfits that land at or under the user’s stated budget on the first pass, with no post-hoc correction needed.	Validates the AI engine and catalog-matching logic are enforcing the hard budget ceiling correctly, not just approximately.
Catalog Swap Rate	The % of recommended items that require an automatic out-of-stock or size-unavailable substitution before display.	A proxy for retailer feed staleness and catalog coverage quality — a high rate signals a partner feed problem, not a styling problem.
Time to Complete an Outfit	The elapsed time from initial input submission to the user selecting or acting on an outfit.	Captures efficiency directly against the core promise: replacing five open tabs with one fast decision.

3. AARRR Funnel Mapping

Funnel Stage	User Action	Metric
Acquisition	User discovers MAISON via Pinterest, Instagram, a styling creator, or a community/content channel.	App/Website Visitors, Sign-Up Conversion Rate, Cost Per Acquisition (CPA)
Activation	User completes their first outfit-recreation session and views the full set of 2–3 curated outfits.	First-Time Activation Rate (%)
Retention	User returns to generate at least one additional outfit after their first session.	7-Day Retention Rate, 30-Day Retention Rate
Revenue	User upgrades to a premium plan for unlimited sessions, saved style history, or priority generation.	Free-to-Paid Conversion Rate, Monthly Recurring Revenue (MRR)
Referral	User shares a recommended outfit or screenshot with a friend before buying.	Referral Rate, K-Factor, Shares Per User

Revenue is intentionally not a primary focus during MVP validation. The initial objective is to validate whether AI-curated, multi-retailer outfits reduce decision fatigue and earn user trust. Monetization is explored only after activation and retention are demonstrated — and,

per the resolution reached during discovery-phase review, only on a revenue model (flat-fee-weighted, commission-normalized) that does not compromise the retailer-agnostic promise (see Section 7).

4. A/B Test Design: Image-First vs. Text-First Capture

4.1 Hypothesis

If the outfit-capture screen leads with a prominent image-upload dropzone rather than a text-prompt field, First-Time Activation Rate will increase by $\geq 15\%$ among users arriving from image-heavy acquisition channels (Pinterest, Instagram), because inspiration-driven shoppers think visually and a text-first flow adds an unnecessary translation step between what they saw and what they want.

4.2 Control Version (A)

The capture screen opens with a text input field (“Describe your occasion and budget”) as the primary element; image upload is available but presented as a secondary, smaller option below it.

4.3 Variant Version (B)

The capture screen opens with a large, primary image-upload dropzone (“Upload your inspiration”); a text prompt is available via a secondary “or describe it instead” link.

4.4 Target Users

50/50 split randomized routing across all newly registered traffic over a 14-day testing cycle, with results segmented post-hoc by acquisition channel to check whether the effect concentrates in image-heavy channels as hypothesized.

4.5 Primary Success Metric

First-Time Activation Rate: % of new users who view a complete set of curated outfits during their first session. The 15% lift target reflects that this is an interface-ordering change, not a new capability — a smaller, more conservative bar than a feature-level test.

4.6 Secondary Metrics

Time to first submission (session latency before upload/text entry) · Drop-off rate before first submission · Budget-field completion rate (whether users bother setting a budget upfront in each variant).

4.7 Success Criteria

- 1. Ship Variant B if: lift $\geq 15\%$, $p < 0.05$, $n \geq 500$ per variant.
- 2. Ship Variant A if: Variant B shows no statistically significant lift, or a negative secondary metric (higher pre-submission drop-off).
- 3. If inconclusive ($p > 0.05$ at end of window): extend the test by 7 days with expanded traffic. If still inconclusive, default to Variant A (text-first), since it serves both personas (Riya and Neha) without assuming visual-first behavior, and flag the result for qualitative interview follow-up split by persona.

5. Dashboard Plan

Dashboard	Metrics Shown	Owner	Cadence
Core Efficacy	Decision Confidence Score · Outfit Completion Rate · Recommendation Acceptance Rate · Budget Adherence Accuracy	Product & AI/Stylist Prompt Engineers	Weekly
Growth & Funnel Health	Unique sign-ups · Onboarding drop-off · First-Time Activation Rate · D7/D30 Retention Cohorts · Referral K-Factor	Growth PM / Marketing	Weekly (first 90 days) → Bi-weekly thereafter
System Stability	Catalog feed latency/staleness · API response latency · Catalog Swap Rate · Buy-link failure rate	Core Engineering	Daily

6. Go-To-Market (GTM) Strategy

Target Audience

Urban women aged 20–30 who discover fashion inspiration on Pinterest or Instagram and experience decision fatigue recreating outfits across retailers (primary persona: Riya), plus busy professionals who need fast, occasion-appropriate outfits from a plain-language brief rather than an image (secondary persona: Neha).

Positioning Statement

For inspiration-driven shoppers trapped in decision fatigue, MAISON is the retailer-agnostic AI stylist that turns a screenshot or a sentence into a complete, budget-ready outfit. Unlike single-retailer apps that only show you their own inventory, MAISON finds the best version of the look — wherever it’s sold.

Launch Channels

1. **Organic Community Seeding:** Pinterest board collaborations and native, non-promotional content inside fashion-forward communities (student groups, budget-styling subreddits and Discords, campus pages).
2. **Short-Form Video Documentation (TikTok/Reels/Shorts):** Raw “screenshot to shopping cart” recordings — a Pinterest save going in, three curated, in-budget outfit cards coming out.
3. **Micro-Influencer Partnerships:** Collaborations with affordable-styling and campus fashion creators to run live “recreate my look, my budget” challenges.

Key Launch Message

“Stop scrolling five apps for one outfit. Show us the look — we’ll find it, in your budget, wherever it’s sold.”

MVP Launch Checklist

- ☐ Hard budget-ceiling validation tested against live, multi-retailer catalog feed
- ☐ Pre-curated fallback outfit sets live for latency-spike or outage degradation
- ☐ Retailer buy-link tracking and attribution verified across all partnered feeds
- ☐ Image encryption in transit and at rest, and deletion-on-request flow tested
- ☐ IP/copyright legal review completed and signed off before the first AI-touching (non-Wizard-of-Oz) user
- ☐ Privacy-compliant anonymous analytics connected
- ☐ AI transparency label and ranking-policy disclosure live on the primary UI
- ☐ Retailer partner onboarding deck and catalog-integration playbook ready for the next partner outreach

7. Pricing & Monetization: The “Retailer-Agnostic, Trust-First” Model

Freemium Tiered Subscription

Free tier — ₹0/month - 3 outfit-recreation sessions per month - Full 2-3 outfit comparison with buy links - Budget, occasion, and size filters - “Rate This Outfit” and “Show Alternatives” feedback

Premium — Individual — ₹299/month - Everything in Free - Unlimited outfit-recreation sessions - Saved style and budget history - Priority generation speed - Early access to wardrobe-completion features (roadmap)

Retailer Partner Program (B2B) - Normalized, disclosed commission rate across all partnered retailers — no pay-to-rank - Priority catalog sync and a fresher feed SLA - Quarterly performance and attribution reports - Featured-partner co-marketing slots, always clearly labeled as sponsored and kept separate from ranked recommendations

Strategic Justification:

Fashion-forward shoppers experience a version of the same guilt cycle seen in other impulse categories: pieces bought in a moment of inspiration that never get styled into a complete look and end up unworn in the closet. A hard paywall on day one would undercut trust in a product whose entire promise is reducing wasted spend, not adding to it — so a genuinely useful free tier proves outfit-completion value before asking for payment.

The commission-normalization commitment on the Retailer Partner Program is not a cosmetic detail — it is the direct resolution to the monetization-vs-positioning conflict identified during discovery-phase review (Discovery & Strategy, Section 17.1): if commission rates were allowed to vary freely by retailer, ranking logic would carry a revenue incentive to favor higher-commission partners over the objectively best match, undermining the “best outfit regardless of retailer” promise this entire product is built on. Normalizing rates and publishing the ranking policy keeps commission revenue supplementary to, rather than in conflict with, that promise.

8. Post-Launch Feedback Plan

Feedback Channels

1. **In-App Micro-Nudge:** A single-tap question surfaced after a buy-link click or session end: “Did this feel like a look you’d actually wear? [Yes, this is me] [Not quite my style].”
2. **Anonymous Community Link:** A link embedded in app settings for users to share feature requests or style-category gaps without linking back to their identity.

Operational Feedback Lifecycle

1. **Collect:** All “Show Alternatives” clicks, low-confidence ratings (1-2), and abandoned sessions are compiled into an automated

pattern log.

2. **Evaluate:** If a specific occasion category (e.g., “wedding guest,” “interview,” “festival”) shows a disproportionately high rate of low-confidence ratings or alternative requests, it flags that styling logic or catalog coverage is weak for that category specifically.
3. **Iterate:** Refine stylist prompt engineering and catalog-matching weights bi-weekly; escalate persistent catalog gaps in a given category to the retailer partnerships team as a feed-expansion priority rather than treating it purely as a prompting problem.