

MAISON — MVP Validation Report

Key Learnings & Roadmap Adjustments, Post-Prototype Review

July 2026

This report documents what using the MAISON prototype as a real user surfaced beyond what the Discovery, PRD, and Metrics documents anticipated on paper. It follows the same structure as those documents so it can sit alongside them as a companion artifact, not a replacement.

1. Executive Summary

Recommendation quality alone was not the whole picture. The prototype validated the core hypothesis — that AI-curated, multi-retailer outfits reduce decision fatigue — but also surfaced a finding the original scope did not account for: users optimize for purchasing convenience, not style fit alone. That finding, and several smaller confirmations of risks already flagged in the Discovery document, are documented below and folded into a revised MVP → V1 → V2 → V3 roadmap.

2. Key Learnings from Validation

Learning 1: Multi-Retailer Checkout Friction — Headline Finding

Problem: MAISON optimizes for outfit quality (style fit + budget), not purchasing convenience. A recommended outfit may require separate checkouts across Myntra, Ajio, Zara, and H&M — meaning multiple orders, multiple deliveries, and multiple return policies for a single outfit.

Evidence: During hands-on testing, users hesitated at the buy-link step whenever an outfit split across three retailers, and several said directly they’d rather accept a slightly less optimal outfit from one retailer than manage separate orders and returns for a marginally better one.

Impact: Users may trade outfit quality for convenience in ways the current ranking (style and budget only) doesn’t account for — meaning a technically optimal recommendation can still lose to a “good enough” one in practice.

Decision: Not solved in the MVP. Style and budget optimization ships first to validate the core hypothesis; convenience becomes an explicit ranking dimension once that core value is proven. Added to the product roadmap as a V2 priority.

V2 Solution — Purchase Optimization Engine: New ranking signals — Retailer Count, Shipping Cost, Delivery Time, and a combined Convenience Score — blended with the existing style/budget fit score into a single ranked list.

Interim, lower-cost option: Before building the full multi-signal engine, ship a simple “Prefer fewer retailers, even if slightly less optimal” toggle as a V1 filter. It validates real appetite for this trade-off before investing in the full V2 ranking rebuild.

Guardrail update required: Once Convenience Score ships, the ranking-transparency disclosure already committed to in the Responsible AI document (“ranked by fit to your brief, style match, and price — not by which retailer pays us”) needs to be updated to also disclose retailer count as a ranking factor. Retailer count isn’t a commission bias, but a convenience-weighted ranking could still quietly favor retailers with the broadest single-source catalogs over smaller partners — a different version of the same fairness concern that disclosure was written to prevent. The fix is the same principle applied consistently: disclose whatever actually drives the ranking.

Learning 2: Retail Catalog Availability Is a Structural Constraint, Not a Data Problem

Problem: The original plan assumed a broad multi-retailer catalog was reachable primarily through affiliate network feeds, with scraping as an informal fallback for gaps.

Evidence: In practice, affiliate feeds cover mainstream fast-fashion categories reasonably well but have real gaps in occasion-specific categories — festive and ethnic wear especially. Scraping was tested informally as a fallback and found unreliable at the granularity outfit generation actually needs: inconsistent product schemas and anti-scraping protections broke matching regularly enough that it isn’t viable as a durable source.

Impact: Catalog coverage quality varies by occasion category, not just by retailer — so “add more retailers” doesn’t fully solve the gap on its own.

Decision: Treat catalog sourcing as a hybrid architecture, not a single sourcing strategy: affiliate feeds for stable, high-volume coverage; prioritized direct partnerships (using the first-mover shortlist criteria already defined in the Discovery document) specifically for gap categories; scraping explicitly ruled out as a scale strategy due to both reliability and likely terms-of-service exposure.

Roadmap placement: Cross-cutting and ongoing, not a single version — tracked against the Catalog Swap Rate KPI already defined in the Metrics & GTM document as the leading indicator of where coverage is weakest.

Learning 3: Cold-Start Personalization

Problem: First-time users have no rating or acceptance history, so early sessions can only be matched against the stated brief (occasion, budget, size, style words), not a personal pattern.

Evidence: Early sessions default to reasonable but generic style-space matching; no meaningful personalization signal exists until a user has rated more than one or two outfits.

Impact: First-session quality reads as “good average match” rather than “tailored to this person” — acceptable for a first session, but a risk if quality doesn’t visibly improve by the second or third.

Decision: MVP relies on explicit onboarding inputs to bootstrap the first session. Behavioral personalization — weighting future recommendations from ratings and accepted or rejected outfits — activates from the second session onward, and matures through the V1 (Ratings, Saves, Feedback) and V2 stages of the roadmap.

Learning 4: Confidence Score as a Noisy Leading Indicator — Confirmed, Not New

Problem: The Discovery document already flagged, before any real usage, that Decision Confidence Score is self-reported and could rise without a matching rise in actual purchase intent.

Evidence: During validation, several sessions produced high confidence ratings (4-5) with no buy-link click at all — the exact pattern anticipated on paper, now observed rather than theoretical.

Impact: Confirms the North Star metric cannot be read in isolation.

Decision: No roadmap change needed here — this is precisely why Decision Confidence Score and Recommendation Acceptance Rate were already specified to be reported together per cohort. This entry is included because confirming a predicted risk is still a real validation finding, distinct from the risk itself.

Learning 5: Budget Flexibility Confirmed in Practice

Problem: Discovery-phase user research (Insight #6) found budget flexibility increases with confidence, which was in tension with the MVP’s hard budget ceiling — a tension the Discovery document deliberately kept as a named trade-off rather than resolving outright.

Evidence: During testing, multiple users asked, unprompted, for some version of “can I see the option that’s a bit over budget but fits better” — before any UI existed to support it.

Impact: The “near-budget, higher-fit” opt-in — previously scoped as an unscheduled Should-Have experiment — has real, observed demand rather than only a theoretical basis.

Decision: Move the near-budget alternative opt-in from an unscheduled Should-Have to an explicit V1 backlog item.

Learning 6: Low-Quality Input Images Degrade Output More Than Expected

Problem: Garment-detection accuracy is highly sensitive to image quality. Cropped screenshots, low resolution, and images with multiple people or outfits in frame reduced isolation accuracy.

Evidence: Several test uploads were compressed social-media re-shares (e.g., Instagram Story screenshots) with visible compression artifacts and partially obscured garments.

Impact: This is a variant of the “Garment Misidentification” failure case already anticipated in the AI Feature document, but validation shows it is driven substantially by input quality, not only model capability.

Decision: Add lightweight upload guidance (“clearer, uncropped photos work best,” with a visual example) and trigger the existing low-confidence indicator more consistently for low-resolution uploads specifically, not only for short or ambiguous text prompts.

Learning 7: Social Sharing Behavior — Confirmed

Problem/Question: Discovery-phase research (Insight #8) hypothesized that users seek external validation by sharing outfit screenshots with friends before buying.

Evidence: During informal validation, users screenshotted outfit result cards to send to friends unprompted, before clicking any buy link.

Impact: Confirms the Referral funnel stage in the Metrics & GTM document reflects real behavior rather than only a funnel-completeness assumption. Since users are already doing this manually, a purpose-built share format is likely to lower friction rather than introduce a new behavior.

Decision: Add a lightweight “Share this outfit” action (a clean, shareable card export) to the V1 backlog.

Learning 8: Wardrobe Memory — Confirmed as Future Scope, No Change

Users have items they already own that could complete or reduce the cost of a recommended outfit. This requires wardrobe-item capture and inventory matching, which is out of scope for MVP or V1. Confirmed as a V3 direction; no roadmap change needed this cycle.

3. Already Shipped — Not New Findings

Two items surfaced during review read like gaps but are already scoped MVP requirements, not new discoveries. Worth being precise about the distinction, since presenting a day-one requirement as a new insight is the one thing that would undercut an otherwise strong validation story:

Item	Where It’s Already Specified	What Actually Changes Going Forward
Product Substitution (find a similar item if one is unavailable)	PRD, User Story 6 — an MVP acceptance criterion	Not a new feature; ongoing focus is execution quality (how visually similar the substitute is), not whether it exists
Explainability (“why this outfit” reasoning)	PRD, Functional Requirement 2; AI Feature document, Section 1	Not a new feature; ongoing focus is reasoning quality and specificity, tracked via low-confidence-rating patterns by occasion category

4. Ranking Transparency Update

The Responsible AI document’s ranking-disclosure line is updated to stay accurate once Convenience Score ships in V2:

Before V2: “Outfits are ranked by fit to your brief, style match, and price — not by which retailer pays us.” **From V2 onward:** “Outfits are ranked by fit to your brief, style match, price, and (where relevant) how many retailers you’d need to order from — never by which retailer pays us.”

This keeps the transparency commitment matched to whatever the ranking logic actually weighs, rather than letting the disclosure go stale the moment a new ranking signal is introduced.

5. Revised Product Roadmap

Stage	Scope
MVP	AI outfit generation · Multi-retailer recommendations · Hard budget-ceiling enforcement · 1-5 confidence rating · Product substitution on stock/size unavailability · “Why this outfit” explanation
V1	Saves · Richer feedback signals · Near-budget, higher-fit alternative (opt-in) · “Share this outfit” export · Weather-aware recommendations · “Prefer fewer retailers” filter (interim convenience signal)
V2	Purchase Optimization Engine (Retailer Count, Shipping Cost, Delivery Time, Convenience Score) ★ · Retailer preference / single-retailer mode · Delivery optimization · Updated ranking-transparency disclosure
V3	Wardrobe AI / Wardrobe Memory · Virtual Try-On · Outfit Calendar · Weather-linked wardrobe planning

Two corrections from the earlier draft roadmap: confidence rating moves from V1 to MVP (it’s a day-one requirement feeding the North Star metric, not an addition), and weather-aware recommendations moves from V3 to V1 (it was already scoped as a Should-Have in the Discovery document’s MoSCoW, not a far-future item alongside virtual try-on).

6. Assumptions That Changed

Initial Assumption	What We Learned	Product Decision
Best outfit = highest style-and-budget fit score	Convenience — retailer count, delivery time, returns — meaningfully changes what users actually choose	Add a Purchase Optimization Engine (V2) with a Convenience Score ranking signal
A broad multi-retailer catalog is reachable primarily through affiliate feeds, with scraping as a fallback	Scraping is unreliable at the granularity outfit generation needs; affiliate feeds have real category gaps, especially festive and ethnic wear	Hybrid catalog architecture — affiliate feeds plus prioritized direct partnerships for gap categories; scraping ruled out as a scale strategy
A hard budget ceiling with no exceptions is the safest default	Users spontaneously ask for near-budget, better-fit alternatives — echoing the flexibility the original user research predicted	Move the near-budget alternative opt-in from an unscheduled Should-Have to an explicit V1 backlog item

7. Assumptions We Got Right From the Start

For balance, and because it matters just as much for an honest account of the process: not everything worth noting is a correction. Two decisions made before any user testing held up exactly as designed, and are worth naming as validated rather than re-labeled as new discoveries:

Original Decision	Status
One recommendation would not be enough — return 2–3 ranked outfits	Correct from day one. Specified in the Discovery document’s MVP scope before any user testing occurred, not a correction made afterward.
AI should assist, not decide — users need visible control and explanation	Correct from day one. Design Principle #5, “Show Alternatives,” and the stylist-persona rationale were founding decisions, shipped as MVP requirements rather than added in response to a walked-back assumption.

8. Framing Note

The most useful thing this report captures isn’t any single fix — it’s that recommendation quality alone wasn’t sufficient, and that users optimize for convenience and control alongside fashion fit, not instead of it. That distinction — between a product that only got more accurate and one that got more usable — is the difference between shipping a feature list and iterating on a real product, and it’s the throughline connecting Learnings 1, 5, 6, and 7 above.