

MAISON — MVP Validation Report

Self-Review Findings & Roadmap Adjustments, Post-Prototype Review

July 2026

This report documents what using the MAISON prototype firsthand surfaced beyond what the Discovery, PRD, and Metrics documents anticipated on paper. It follows the same structure as those documents so it can sit alongside them as a companion artifact, not a replacement.

1. Executive Summary

Recommendation quality alone was not the whole picture. Reviewing the live prototype directly — and reasoning through it against the Riya and Neha personas and the original Discovery-phase research — surfaced a finding the original scope did not account for: users are likely to optimize for purchasing convenience, not style fit alone. That finding, alongside several places where the live product re-surfaces tensions the Discovery document had already flagged on paper, is documented below and folded into a revised MVP → V1 → V2 → V3 roadmap.

2. Methodology

This round of validation was a **structured self-review of the live prototype**, not a moderated usability study with external participants. It consisted of: using the deployed prototype directly across representative inputs (image uploads, text prompts, varying budgets and occasions); reasoning through each interaction against the two validated personas (Riya and Neha) rather than assuming a single “average user” reaction; and cross-checking observed product behavior against the risks and open questions the Discovery and PRD documents had already flagged before any prototype existed.

This is a meaningful method for catching design and logic gaps early, and it’s the right next step after a documentation-only phase — but it is not a substitute for testing with real, external users. Findings below are labeled accordingly: some are things directly observable in the product as built; others are reasoned judgments or hypotheses carried forward from the original research that this review reinforced but did not independently confirm with other people. Section 9 proposes the next validation step to close that gap.

3. Key Learnings from Review

Learning 1: Multi-Retailer Checkout Friction — Headline Finding

Directly observable in the product as built.

Problem: MAISON optimizes for outfit quality (style fit + budget), not purchasing convenience. A recommended outfit may require separate checkouts across Myntra, Ajio, Zara, and H&M — meaning multiple orders, multiple deliveries, and multiple return policies for a single outfit.

Observation: Walking through the buy-link step directly, an outfit split across three retailers visibly multiplies the steps needed to actually complete a purchase. Reasoning this through against the personas: Neha (time-constrained, wants one-session completion) is the more likely of the two to abandon here rather than manage three separate orders for one outfit.

Impact: Users may trade outfit quality for convenience in ways the current ranking (style and budget only) doesn’t account for — a technically optimal recommendation can still lose to a “good enough” one in practice.

Decision: Not solved in the MVP. Style and budget optimization ships first to validate the core hypothesis; convenience becomes an explicit ranking dimension once that core value is proven with real users. Added to the product roadmap as a V2 priority.

V2 Solution — Purchase Optimization Engine: New ranking signals — Retailer Count, Shipping Cost, Delivery Time, and a combined Convenience Score — blended with the existing style/budget fit score into a single ranked list.

Interim, lower-cost option: Before building the full multi-signal engine, ship a simple “Prefer fewer retailers, even if slightly less optimal” toggle as a V1 filter. It’s a cheap way to test real appetite for this trade-off with actual users before investing in the full V2 ranking rebuild.

Guardrail update required: Once Convenience Score ships, the ranking-transparency disclosure already committed to in the Responsible AI document (“ranked by fit to your brief, style match, and price — not by which retailer pays us”) needs to be updated to also disclose retailer count as a ranking factor. Retailer count isn’t a commission bias, but a convenience-weighted ranking could still quietly favor retailers with the broadest single-source catalogs over smaller partners — a different version of the same fairness concern that disclosure was written to prevent.

Learning 2: Retail Catalog Availability Is a Structural Constraint, Not a Data Problem

Reasoned judgment, carried forward from Discovery-phase analysis and reinforced by reviewing the prototype’s demo data.

Problem: The original plan assumed a broad multi-retailer catalog was reachable primarily through affiliate network feeds, with scraping as an informal fallback for gaps.

Reasoning: This was already anticipated as a risk in the Discovery document (Section 4’s incentive-gap analysis, and the catalog-access risk in Section 15). Reviewing how thin the demo/prototype catalog coverage is for occasion-specific categories — festive and ethnic wear especially — reinforces that concern rather than introducing a new one. This has not yet been tested against a real, live retailer feed or a real scraping attempt; it remains a structural risk based on how affiliate networks and retailer APIs are known to work, not an empirically measured failure.

Impact: Catalog coverage quality is likely to vary by occasion category, not just by retailer — so “add more retailers” probably doesn’t fully solve the gap on its own.

Decision: Treat catalog sourcing as a hybrid architecture, not a single sourcing strategy: affiliate feeds for stable, high-volume coverage; prioritized direct partnerships (using the first-mover shortlist criteria already defined in the Discovery document) specifically for gap categories; scraping treated with caution as a scale strategy pending an actual reliability and terms-of-service assessment.

Roadmap placement: Cross-cutting and ongoing — tracked against the Catalog Swap Rate KPI already defined in the Metrics & GTM document once real retailer feeds are live, since that’s the metric that will actually confirm or disconfirm this concern.

Learning 3: Cold-Start Personalization

Structural/design reasoning — a known limitation of the current architecture, not something requiring user testing to identify.

Problem: First-time users have no rating or acceptance history, so early sessions can only be matched against the stated brief (occasion, budget, size, style words), not a personal pattern.

Reasoning: This follows directly from how the personalization engine is scoped in the AI Feature document — there is no behavioral signal until a user has rated more than one or two outfits, by design.

Impact: First-session quality will read as “good average match” rather than “tailored to this person” — acceptable for a first session, but worth monitoring once real usage data exists to see whether quality visibly improves by the second or third session.

Decision: MVP relies on explicit onboarding inputs to bootstrap the first session. Behavioral personalization activates from the second session onward, maturing through the V1 (Ratings, Saves, Feedback) and V2 stages of the roadmap. No change to this plan; flagged here as a known limitation worth watching once real usage data is available.

Learning 4: Confidence Score as a Noisy Leading Indicator — Anticipated Risk, Not Yet Testable

Cannot be confirmed from this round of review — the current prototype does not have persistent, multi-session analytics.

Problem: The Discovery document already flagged, before any prototype existed, that Decision Confidence Score is self-reported and could rise without a matching rise in actual purchase intent.

Status: This review could not confirm or disconfirm this pattern, since the prototype doesn’t yet log confidence ratings against buy-link clicks across real sessions. It remains exactly what it was in the Discovery document: a predicted risk, not an observed one. It’s listed here as a reminder that this needs real session data to validate, not as a new finding.

Decision: No roadmap change. Decision Confidence Score and Recommendation Acceptance Rate are already specified to be reported together per cohort (Metrics & GTM document, Section 5) — this entry flags that the analytics needed to actually populate that dashboard aren’t live yet in the prototype, which is worth tracking as a build dependency.

Learning 5: Budget Flexibility — Reinforced by Firsthand Use

Directly observable through first-person use of the prototype.

Problem: Discovery-phase user research (Insight #6) found budget flexibility increases with confidence, which was in tension with the MVP’s hard budget ceiling — a tension the Discovery document deliberately kept as a named trade-off rather than resolving outright.

Observation: Testing the budget-ceiling behavior directly surfaced the same tension: seeing a well-matched outfit excluded from results for being marginally over budget prompts the same “let me see it anyway” reaction the original research anticipated. This is a first-person reaction during review, not a finding from multiple independent users, but it’s a concrete enough experience of the friction to move the response up in priority.

Impact: The “near-budget, higher-fit” opt-in — previously scoped as an unscheduled Should-Have — is worth prioritizing sooner, though confirming real demand (versus one reviewer’s reaction) is still worth doing with actual users before committing significant build effort.

Decision: Move the near-budget alternative opt-in from an unscheduled Should-Have to an explicit V1 backlog item, framed as a hypothesis to validate with real usage once built, not a confirmed requirement.

Learning 6: Low-Quality Input Images Degrade Output More Than Expected

Directly observable — tested firsthand with a range of representative images.

Problem: Garment-detection accuracy is sensitive to image quality. Cropped screenshots, low resolution, and images with multiple people or outfits in frame reduce isolation accuracy.

Observation: Uploading a handful of representative screenshots — including compressed social-media re-shares with visible compression artifacts and partially obscured garments — showed a visible drop in output quality compared to clean, high-resolution inspiration images.

Impact: This is a variant of the “Garment Misidentification” failure case already anticipated in the AI Feature document, now observed directly rather than only theorized. It’s a small, self-tested sample, not a systematic accuracy benchmark, but the direction is clear enough to act on.

Decision: Add lightweight upload guidance (“clearer, uncropped photos work best,” with a visual example) and trigger the existing low-confidence indicator more consistently for low-resolution uploads specifically, not only for short or ambiguous text prompts.

Learning 7: Social Sharing Behavior — Hypothesis Carried Forward, Not Yet Re-Validated

Not confirmed by this review — no external users were involved in this round.

Problem/Question: Discovery-phase research (Insight #8) hypothesized that users seek external validation by sharing outfit screenshots with friends before buying.

Status: This review did not involve other users, so this behavior could not be observed or confirmed here. It remains a carried-forward hypothesis from the original research, not a new finding from this round. It’s included in this report because the outfit-card format doesn’t currently have a dedicated share action, and if Insight #8 holds, building one deliberately would lower friction on behavior users are otherwise likely to do manually (via screenshot).

Decision: Add a lightweight “Share this outfit” action to the V1 backlog as a low-cost bet consistent with the original research — but treat it as speculative until confirmed with real users, not as a validated finding.

Learning 8: Wardrobe Memory — Confirmed as Future Scope, No Change

Users have items they already own that could complete or reduce the cost of a recommended outfit. This requires wardrobe-item capture and inventory matching, which is out of scope for MVP or V1. Confirmed as a V3 direction based on the original roadmap reasoning; no new information from this review.

4. Already Shipped — Not New Findings

Two items surfaced during review read like gaps but are already scoped MVP requirements, not new discoveries. Worth being precise about the distinction, since presenting a day-one requirement as a new insight would undercut an otherwise honest validation story:

Item	Where It’s Already Specified	What Actually Changes Going Forward
Product Substitution (find a similar item if one is unavailable)	PRD, User Story 6 — an MVP acceptance criterion	Not a new feature; ongoing focus is execution quality (how visually similar the substitute is), not whether it exists
Explainability (“why this outfit” reasoning)	PRD, Functional Requirement 2; AI Feature document, Section 1	Not a new feature; ongoing focus is reasoning quality and specificity, tracked via low-confidence-rating patterns by occasion category once real usage data exists

5. Ranking Transparency Update

The Responsible AI document’s ranking-disclosure line is updated to stay accurate once Convenience Score ships in V2:

Before V2: “Outfits are ranked by fit to your brief, style match, and price — not by which retailer pays us.” **From V2 onward:** “Outfits are ranked by fit to your brief, style match, price, and (where relevant) how many retailers you’d need to order from — never by which retailer pays us.”

This keeps the transparency commitment matched to whatever the ranking logic actually weighs, rather than letting the disclosure go stale the moment a new ranking signal is introduced.

6. Revised Product Roadmap

Stage	Scope
MVP	AI outfit generation · Multi-retailer recommendations · Hard budget-ceiling enforcement · 1–5 confidence rating · Product substitution on stock/size unavailability · “Why this outfit” explanation
V1	Saves · Richer feedback signals · Near-budget, higher-fit alternative (opt-in, to

	be validated) · “Share this outfit” export (speculative, pending validation) · Weather-aware recommendations · “Prefer fewer retailers” filter (interim convenience signal)
V2	Purchase Optimization Engine (Retailer Count, Shipping Cost, Delivery Time, Convenience Score) ★ · Retailer preference / single-retailer mode · Delivery optimization · Updated ranking-transparency disclosure
V3	Wardrobe AI / Wardrobe Memory · Virtual Try-On · Outfit Calendar · Weather-linked wardrobe planning

Two corrections from the earlier draft roadmap: confidence rating moves from V1 to MVP (it’s a day-one requirement feeding the North Star metric, not an addition), and weather-aware recommendations moves from V3 to V1 (it was already scoped as a Should-Have in the Discovery document’s MoSCoW, not a far-future item alongside virtual try-on).

7. Assumptions That Changed

Initial Assumption	What This Review Surfaced	Product Decision
Best outfit = highest style-and-budget fit score	Walking through the checkout flow directly shows convenience (retailer count, delivery, returns) is a real trade-off users would plausibly weigh, not only style fit	Add a Purchase Optimization Engine (V2) with a Convenience Score ranking signal, validated with real users before full build
A broad multi-retailer catalog is reachable primarily through affiliate feeds, with scraping as a fallback	Reasoning through the catalog architecture against known affiliate-network limitations reinforces that category-specific gaps (festive/ethnic wear especially) are a structural risk, not a scraping-solvable one	Hybrid catalog architecture — affiliate feeds plus prioritized direct partnerships for gap categories; scraping treated cautiously pending a real reliability assessment
A hard budget ceiling with no exceptions is the safest default	Firsthand use of the budget-ceiling behavior reproduced the exact tension Insight #6 predicted from the original research	Move the near-budget alternative opt-in from an unscheduled idea to an explicit V1 backlog item, to be validated with real users

8. Assumptions We Got Right From the Start

For balance: not everything worth noting is a correction. Two decisions made before any prototype existed held up exactly as designed, and are worth naming as validated rather than re-labeled as new discoveries:

Original Decision	Status
One recommendation would not be enough — return 2-3 ranked outfits	Correct from day one. Specified in the Discovery document’s MVP scope before any prototype existed, not a correction made afterward.
AI should assist, not decide — users need visible control and explanation	Correct from day one. Design Principle #5, “Show Alternatives,” and the stylist-persona rationale were founding decisions, shipped as MVP requirements rather than added in response to a walked-back assumption.

9. Recommended Next Validation Step

This review sharpened the roadmap and caught real design gaps, but three of the findings above (Learnings 4, 5, and 7) are reasoned judgments or single-reviewer reactions, not confirmed patterns across real users. Before committing significant build effort to the V1/V2 items motivated by them — particularly the Purchase Optimization Engine and the share-outfit feature — the next step should be a small, structured round of testing with 5-8 real users from the Riya/Neha personas, using the live prototype, specifically checking:

- 1. Do independent users hit the same multi-retailer checkout hesitation observed here (Learning 1)?
- 2. Do independent users spontaneously request near-budget alternatives (Learning 5), or was that a single reviewer’s reaction?
- 3. Do independent users share outfit results with others unprompted (Learning 7), confirming Insight #8 in this specific product context?

This keeps the roadmap decisions in this report honestly labeled as *informed by* self-review, while giving a clear, cheap next step to convert them into *validated by users* before the V2 investment is made.

10. Framing Note

The most useful thing this report captures isn’t any single fix — it’s the discipline of separating what was actually observed in the product from what was reasoned through against personas and prior research, and being honest about which is which. That distinction is what turns a self-review into something a real user-validation round can build on, rather than a set of claims that would not hold up under a follow-up question about methodology.